

# CS324: Deep Learning

## Assignment 1

From Jianguo Zhang

February 20, 2022

### 1 Part I: the perceptron (20 points)

In this first task you're asked to implement and test a simple artificial neuron: a perceptron (see **perceptron-slides.pdf**).

#### 1.1 Task 1

Generate a dataset of points in  $\mathbb{R}^2$ . To do this, define two Gaussian distributions and sample 200 points from each. Your dataset should then contain a total of 400 points, 200 from each distribution. Keep 160 points per distribution as the training (320 in total), 40 for the test (80 in total).

#### 1.2 Task 2

Implement the perceptron following the specs in **perceptron.py** and the pseudocode in **perceptronslides.pdf**.

#### 1.3 Task 3

Train the perceptron on the training data (320 points) and test it on the remaining 80 test points. Compute the classification accuracy on the test set.

#### 1.4 Task 4

Experiment with different sets of points (generated as described in Task 1). What happens during the training if the means of the two Gaussians are too close and/or if their variance is too high?

### 2 Part II: the multi-layer perceptron (65 points)

In this second part of Assignment I you're asked to implement a multi-layer perceptron using numpy. Using scikit-learn and the **make\_moons** method<sup>1</sup>, create a dataset of 2,000 two-dimensional points. Let  $S$  denote the dataset, i.e., the set of tuples  $\{(x^{(0),s}, t^s)\}_{s=1}^S$ , where  $x^{(0),s}$  is the  $s$ -th element of the dataset and  $t^s$  is its label. Further let  $d_0$  be the dimension of the input space and  $d_n$  the dimension of the output space. In this assignment we want the labels to be one-hot encoded<sup>2</sup>. The network you will build will have  $N$  layers (including the output layer). In particular, the structure will be as follows:

- Each layer  $l = 1, \dots, N$  first applies the affine mapping

$$\tilde{x}^{(l)} = W^{(l)}x^{(l-1)} + b^{(l)},$$

---

<sup>1</sup>[https://scikit-learn.org/stable/modules/generated/sklearn.datasets.make\\_moons.html#sklearn.datasets.make\\_moons](https://scikit-learn.org/stable/modules/generated/sklearn.datasets.make_moons.html#sklearn.datasets.make_moons)

<sup>2</sup>Remember to transform the original dataset labels using one-hot encoding <https://en.wikipedia.org/wiki/One-hot>

where  $W^{(l)} \in \mathbb{R}^{d_l \times d_{(l-1)}}$  is the matrix of the weight parameters and  $b^{(l)} \in \mathbb{R}^{d_l}$  is the vector of biases. Given  $\tilde{x}^{(l)}$ , the activation of the  $l$ -th layer is computed using a ReLU unit

$$x^{(l)} = \max(0, \tilde{x}^{(l)}).$$

- The output layer (i.e., the  $N$ -th layer) first applies the affine mapping

$$\tilde{x}^{(N)} = W^{(N)}x^{(N-1)} + b^{(N)},$$

and then uses the softmax activation function (instead of the ReLU of the previous layers) to compute a valid probability mass function (pmf)

$$x^{(N)} = \text{softmax}(\tilde{x}^{(N)}) = \frac{\exp(\tilde{x}^{(N)})}{\sum_{i=1}^{d_N} \exp(\tilde{x}^{(N)})_i}.$$

Note that both max and exp are element-wise operations.

- Finally, compute the cross entropy loss  $L$  between the predicted and the actual label,

$$L(x^{(N)}, t) = - \sum_i t_i \log x_i^{(N)}.$$

## 2.1 Task 1

Implement the MLP architecture by completing the files **mlp\_numpy.py** and **modules.py**.

## 2.2 Task 2

Implement training and testing script in **train\_mlp\_numpy.py**. (Please keep 70% of the dataset for training and the remaining 30% for testing. Note that this is a random split of 70% and 30% )

## 2.3 Task 3

Using the default values of the parameters, report the results of your experiments using a jupyter notebook where you show the accuracy curves for both training and test data.

# 3 Part III: stochastic gradient descent (15 points)

In this third part of Assignment I, you will implement an alternative training method in **train\_mlp\_numpy.py** based on stochastic gradient descent.

## 3.1 Task 1

Modify the train method in **train\_mlp\_numpy.py** to accept a parameter that allows the user to specify if the training has to be performed using batch gradient descent (which you should have implemented in Part II) or stochastic gradient descent.

## 3.2 Task 2

Using the default values of the parameters, report the results of your experiments using a jupyter notebook where you show the accuracy curves for both training and test data.

## 4 Submission instructions

The submission will include:

- A written report describing what you did, the results and your analysis.
- Code for producing **all** results for all parts and tasks.
- Instructions on how to run the code.

Create a ZIP archive with the submission of Assignment 1 (all parts and tasks). Give the ZIP file the name **studentnumber\_assignment1.zip**, where you insert your student number. Please submit the archive through the Blackboard.

Make sure all files needed to run your code are included or you may be given 0 points for it.

**The deadline for assignment 1 (all parts and all tasks) is the 21st of March 2022 at 23:55 (Beijing Time).**